

# CA2 Individual Report

Name	Ong Jin Kai
Student Id	P2429465
Class	DIT/FT/2B/21
Github Repository URL	<a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai</a>
Github Account ID	cutiepatootiekai

For each criterion, provide links to pull requests/commits/files that demonstrate the completion of the requirement. Replace each “?” with your Self Rating.

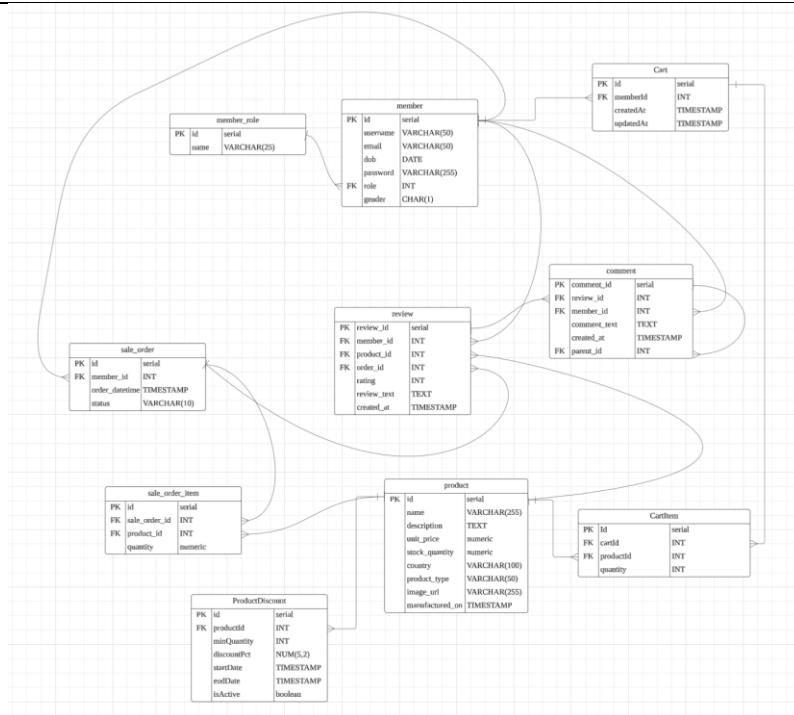
For Self Rating, you may rate yourself accordingly if you feel that you:

- 0 Have little or **no** understanding. and did not attempt the requirement.
- 1 Have **limited** understanding to demonstrate competency for the criterion.
- 2 Have **basic** understanding and only able to replicate examples from tutorials/practicals.
- 3 Have **adequate** understanding and can extend from what you have learned to fulfil specifications.
- 4 Have **solid** understanding in the specific criterion, able work on the requirement without much references.
- 5 Have **excellent** understanding and implemented the requirement according to latest industry guidelines, best practices and documentations.

## Important

- a) You are required to provide for each criterion:
  - **Documentation** and description of the work done.
  - **One to three** of your best implementations with URL **link** to respective repository files/commits/pull requests.
  - You should also provide **screenshots** where relevant.
- b) You are to ensure the hyperlink in this document works. **Failure to do so will result in a 50% deduction of marks.**

No	Criterion	Describe What Was Done	Self Rating
1	Database Design & ORM Modeling		3/5



The new models are Cart, CartItem, ProductDiscount.

<https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/prisma/schema.prisma>

#### Stock only reduced at checkout

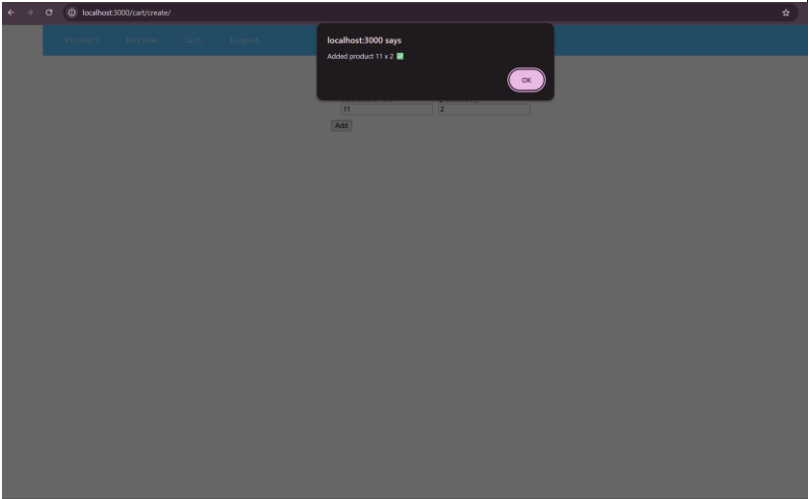
- We don't "reserve" items when they're added to the cart.
- Whoever checks out first gets the stock; if stock isn't enough, the order fails at checkout.
- This keeps things simple—no timers or cleanup for holds.

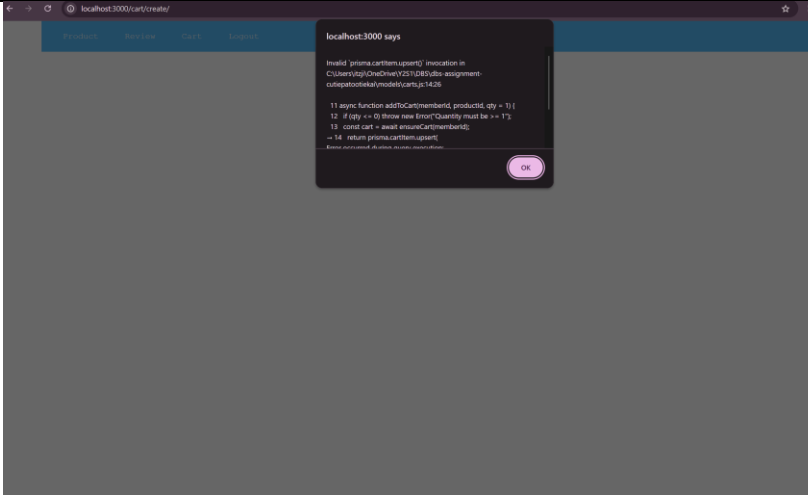
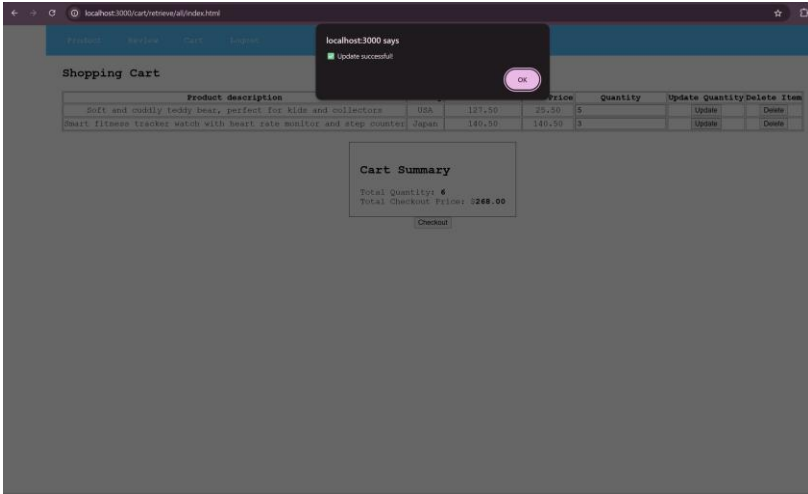
#### One best discount per product line

- If multiple discounts could apply, we pick the **single** one that gives the **lowest price**.
- We **don't** stack discounts.

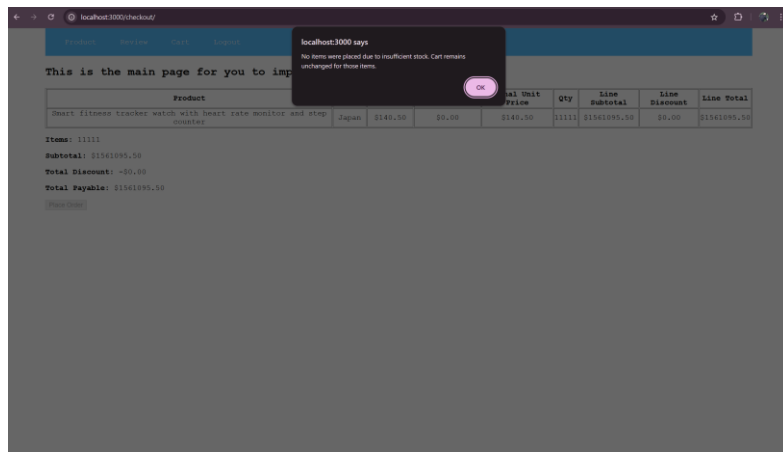
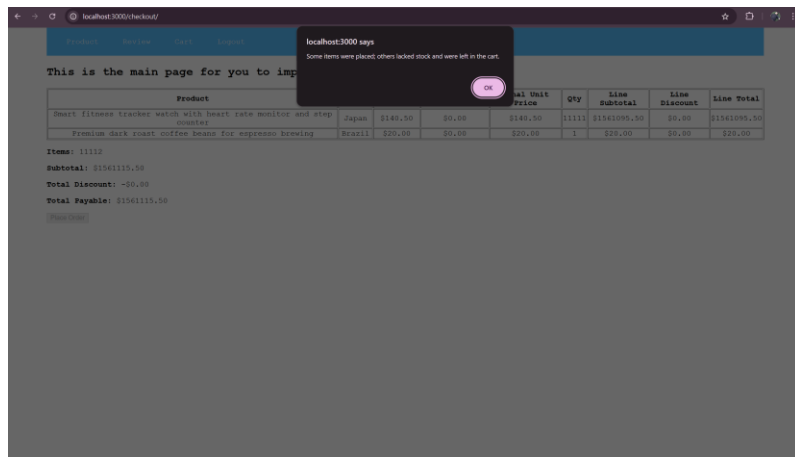
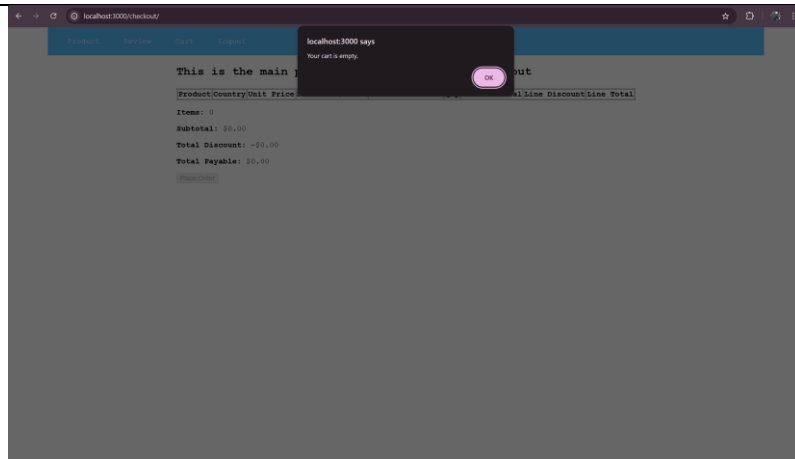
#### When a discount is valid

- It must be **active**,
- Today must be **between start and end dates** (no end date = still okay),
- And you must buy at least the **minimum quantity** (if any).
- That's how most promos work in shops, and it's easy to check in SQL.

	Implementati on	<p><b><u>Backend</u></b></p> <ul style="list-style-type: none"><li>• createCartItems</li><li>• updateCartItems</li><li>• retrieveCartItems</li><li>• deleteCartItems</li><li>• getCartSummary</li><li>• ensureCart</li><li>• addToCart</li><li>• updateLine</li><li>• removeLine</li><li>• getCart</li><li>• getSummary</li></ul> <p><b><u>Frontend</u></b></p> <ul style="list-style-type: none"><li>• refreshCart</li><li>• fetchCartItems</li><li>• fetchCartSummary</li></ul> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/controllers/cartsController.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/controllers/cartsController.js</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/carts.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/carts.js</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/cart">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/cart</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/routes/carts.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/routes/carts.js</a></p> 	
--	--------------------	--	--

		 	
3	Checkout feature Implementation	<p><b>Backend</b></p> <ul style="list-style-type: none"> <li>• <i>getCheckoutSummary</i></li> <li>• <i>bestDiscountPerUnit</i></li> <li>• <i>computeSummary</i></li> </ul> <p><b>Frontend</b></p> <ul style="list-style-type: none"> <li>• <i>currency</i></li> <li>• <i>renderSummary</i></li> <li>• <i>loadSummary</i></li> </ul> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/checkout">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/checkout</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/controllers/checkoutController.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/controllers/checkoutController.js</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/checkout.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/checkout.js</a></p> <p><a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/routes/checkout.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/routes/checkout.js</a></p>	4/5

		<div><div><div><div>Product</div><div>Review</div><div>Cart</div><div>Logout</div></div><div><div>This is the main page for you to implement checkout</div><table><tr><th>Product</th><th>Country</th><th>Unit Price</th><th>Discount</th><th>Unit</th><th>Final Unit Price</th><th>Qty</th><th>Line Subtotal</th><th>Line Discount</th><th>Line Total</th></tr><tr><td>Soft and cuddly teddy bear, perfect for kids and collectors</td><td>USA</td><td>\$25.50</td><td>-2.55</td><td></td><td>\$22.95</td><td>3</td><td>\$76.50</td><td>-7.65</td><td>\$68.85</td></tr></table><div>Items: 3</div><div>Subtotal: \$76.50</div><div>Total Discount: -7.65</div><div>Total Payable: \$68.85</div><div>Place Order</div></div></div></div> <div><p>Fig 1: When product quantity is 3, discount is 10%</p></div> <div><div><div><div>Product</div><div>Review</div><div>Cart</div><div>Logout</div></div><div><div>This is the main page for you to implement checkout</div><table><tr><th>Product</th><th>Country</th><th>Unit Price</th><th>Discount</th><th>Unit</th><th>Final Unit Price</th><th>Qty</th><th>Line Subtotal</th><th>Line Discount</th><th>Line Total</th></tr><tr><td>Soft and cuddly teddy bear, perfect for kids and collectors</td><td>USA</td><td>\$25.50</td><td>-3.82</td><td></td><td>\$21.68</td><td>5</td><td>\$127.50</td><td>-19.13</td><td>\$108.38</td></tr></table><div>Items: 5</div><div>Subtotal: \$127.50</div><div>Total Discount: -19.13</div><div>Total Payable: \$108.38</div><div>Place Order</div></div></div></div> <div><p>Fig 2: When product quantity is 5, discount is 15%</p></div> <div><p>Based on the screenshots above, it shows the extensibility of the feature as if there are 2 different discounts for the same product, the feature will select the 1 with the better discounts.</p></div>	Product	Country	Unit Price	Discount	Unit	Final Unit Price	Qty	Line Subtotal	Line Discount	Line Total	Soft and cuddly teddy bear, perfect for kids and collectors	USA	\$25.50	-2.55		\$22.95	3	\$76.50	-7.65	\$68.85	Product	Country	Unit Price	Discount	Unit	Final Unit Price	Qty	Line Subtotal	Line Discount	Line Total	Soft and cuddly teddy bear, perfect for kids and collectors	USA	\$25.50	-3.82		\$21.68	5	\$127.50	-19.13	\$108.38	
Product	Country	Unit Price	Discount	Unit	Final Unit Price	Qty	Line Subtotal	Line Discount	Line Total																																		
Soft and cuddly teddy bear, perfect for kids and collectors	USA	\$25.50	-2.55		\$22.95	3	\$76.50	-7.65	\$68.85																																		
Product	Country	Unit Price	Discount	Unit	Final Unit Price	Qty	Line Subtotal	Line Discount	Line Total																																		
Soft and cuddly teddy bear, perfect for kids and collectors	USA	\$25.50	-3.82		\$21.68	5	\$127.50	-19.13	\$108.38																																		
4	Transaction Management	<div><p>My stored procedure is <code>place_orders(IN p_member_id integer)</code></p><div><div><div><div>Product</div><div>Review</div><div>Cart</div><div>Logout</div></div><div><div>This is the main page for you to im</div><table><tr><th>Product</th><th>Unit Price</th><th>Qty</th><th>Line Subtotal</th><th>Line Discount</th><th>Line Total</th></tr><tr><td>Soft and cuddly teddy bear, perfect for kids and co</td><td>\$21.68</td><td>5</td><td>\$127.50</td><td>-19.13</td><td>\$108.38</td></tr></table><div>Items: 5</div><div>Subtotal: \$127.50</div><div>Total Discount: -19.13</div><div>Total Payable: \$108.38</div><div>Place Order</div></div></div><div>localhost:3000 says Order placement completed. All items were processed successfully.</div></div></div> <div></div>	Product	Unit Price	Qty	Line Subtotal	Line Discount	Line Total	Soft and cuddly teddy bear, perfect for kids and co	\$21.68	5	\$127.50	-19.13	\$108.38	3/5																												
Product	Unit Price	Qty	Line Subtotal	Line Discount	Line Total																																						
Soft and cuddly teddy bear, perfect for kids and co	\$21.68	5	\$127.50	-19.13	\$108.38																																						



### Backend

- *confirmCheckout*
- *confirm*
- *placeOrdersForMember*

<https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/routes/checkout.js>

<https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/controllers/checkoutController.js>

		<a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/checkout.js">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/blob/master/models/checkout.js</a>  <a href="https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/checkout">https://github.com/soc-DBS/dbs-assignment-cutiepatootiekai/tree/master/public/checkout</a>	
5	Indexing	<p><b>1. Case-insensitive firm lookup</b></p> <pre>SELECT id, firm_name FROM manufacturer WHERE lower(firm_name) = lower('Acme Holdings');</pre> <p>Purpose: Find a company by name regardless of letter case. Index: B-tree function(express) index</p> <pre>CREATE INDEX idx_lower_firm_name ON manufacturer (lower(firm_name));</pre> <p><b>Before:</b></p> <div> <p><b>QUERY PLAN</b></p> <p>text</p> <p>Gather (cost=1000.00..26571.12 rows=4091 width=23) (actual time=250.103..258.638 rows=0 loops=1)</p> <p>Workers Planned: 2</p> <p>Workers Launched: 2</p> <p>Buffers: shared hit=15994 read=4054</p> <p>-&gt; Parallel Seq Scan on manufacturer (cost=0.00..25162.03 rows=1705 width=23) (actual time=206.880..206.880 rows=0 loop=1)</p> <p>Filter: (lower((firm_name)::text) = 'acme holdings':text)</p> <p>Rows Removed by Filter: 272748</p> <p>Buffers: shared hit=15994 read=4054</p> <p>Planning:</p> <p>Buffers: shared hit=156 read=7</p> <p>Planning Time: 3.455 ms</p> <p>Execution Time: 258.662 ms</p> </div> <p><b>After:</b></p> <div> <p><b>QUERY PLAN</b></p> <p>text</p> <p>Bitmap Heap Scan on manufacturer (cost=88.13..10207.77 rows=4091 width=23) (actual time=0.059..0.059 rows=0 loops=1)</p> <p>Recheck Cond: (lower((firm_name)::text) = 'acme holdings':text)</p> <p>Buffers: shared read=3</p> <p>-&gt; Bitmap Index Scan on idx_lower_firm_name (cost=0.00..87.11 rows=4091 width=0) (actual time=0.058..0.058 rows=0 loop=1)</p> <p>Index Cond: (lower((firm_name)::text) = 'acme holdings':text)</p> <p>Buffers: shared read=3</p> <p>Planning:</p> <p>Buffers: shared hit=21 read=1</p> <p>Planning Time: 1.538 ms</p> <p>Execution Time: 0.085 ms</p> </div> <p>Reason: Creating a B-tree expression index on lower(firm_name) made the predicate indexable, so the optimizer replaced a Sequential Scan with a Bitmap Index Scan + Bitmap Heap Scan, avoiding a full table read and per-row LOWER() evaluation.</p>	3/5

**2. Recent companies by founded date (latest first)**

```
SELECT id, firm_name, founded_since
FROM manufacturer
WHERE founded_since >= DATE '2020-01-01'
ORDER BY founded_since DESC
LIMIT 100;
```

Purpose: List recently founded firms, newest first, top 100.  
Index: B-tree Partial, covering index

```
CREATE INDEX idx_recent_founded_cover
ON manufacturer (founded_since DESC)
INCLUDE (id, firm_name)
WHERE founded_since >= DATE '2020-01-01';
```

***Before:***

QUERY PLAN
text
Limit (cost=25847.46..25859.12 rows=100 width=27) (actual time=121.724..139.832 rows=100 loops=1)
Buffers: shared hit=16070 read=4052
-> Gather Merge (cost=25847.46..29130.69 rows=28140 width=27) (actual time=121.722..139.822 rows=100 loops=1)
Workers Planned: 2
Workers Launched: 2
Buffers: shared hit=16070 read=4052
-> Sort (cost=24847.43..24882.61 rows=14070 width=27) (actual time=79.041..79.046 rows=84 loops=3)
Sort Key: founded_since DESC
Sort Method: top-N heapsort Memory: 36kB
Buffers: shared hit=16070 read=4052
Worker 0: Sort Method: top-N heapsort Memory: 37kB
Worker 1: Sort Method: top-N heapsort Memory: 37kB
-> Parallel Seq Scan on manufacturer (cost=0.00..24309.69 rows=14070 width=27) (actual time=0.188..77.016 rows=1086...
Filter: (founded_since >= '2020-01-01'::date)
Rows Removed by Filter: 261880
Buffers: shared hit=15996 read=4052
Planning:
Buffers: shared hit=8
Planning Time: 0.796 ms
Execution Time: 139.903 ms

***After:***

QUERY PLAN
text
Limit (cost=0.29..4.27 rows=100 width=27) (actual time=0.366..0.382 rows=100 loops=1)
Buffers: shared hit=1 read=2
-> Index Only Scan using idx_recent_founded_cover on manufacturer (cost=0.29..1344.97 rows=33767 width=27) (actual time=0.365..0.373 rows=100 loop...
Heap Fetches: 0
Buffers: shared hit=1 read=2
Planning:
Buffers: shared hit=24 read=1
Planning Time: 1.647 ms
Execution Time: 0.397 ms

Reason : Creating a partial, covering B-tree on (founded\_since DESC)  
INCLUDE (id, firm\_name) WHERE founded\_since >= '2020-01-01'  
aligned the WHERE and ORDER BY, enabling an Index Only Scan with  
early LIMIT and removing both the full scan and the sort.



### 3. Country filter (USA) sorted by name

```
SELECT id, firm_name, origin
FROM manufacturer
WHERE origin = 'USA'
ORDER BY firm_name
LIMIT 20;
```

Purpose: First 20 USA firms, alphabetically.  
Index: B-tree composite index

```
CREATE INDEX idx_manu_origin_name ON manufacturer (origin,
firm_name);
```

*Before:*

QUERY PLAN	
text	
1	Limit (cost=27102.67..27105.00 rows=20 width=30) (actual time=124.338..135.312 rows=20 loops=1)
2	Buffers: shared hit=15515 read=4607
3	-> Gather Merge (cost=27102.67..42825.77 rows=134760 width=30) (actual time=124.336..135.307 rows=20 loops=1)
4	Workers Planned: 2
5	Workers Launched: 2
6	Buffers: shared hit=15515 read=4607
7	-> Sort (cost=26102.65..26271.10 rows=67380 width=30) (actual time=81.233..81.234 rows=17 loops=3)
8	Sort Key: firm_name
9	Sort Method: top-N heapsort Memory: 27kB
10	Buffers: shared hit=15515 read=4607
11	Worker 0: Sort Method: top-N heapsort Memory: 27kB
12	Worker 1: Sort Method: top-N heapsort Memory: 27kB
13	-> Parallel Seq Scan on manufacturer (cost=0.00..24309.69 rows=67380 width=30) (actual time=0.091..59.949 rows=5438..)
14	Filter: ((origin)::text = 'USA')::text
15	Rows Removed by Filter: 218365
16	Buffers: shared hit=15441 read=4607
17	Planning:
18	Buffers: shared hit=6 dirtied=1
19	Planning Time: 0.964 ms
20	Execution Time: 135.386 ms

*After:*

QUERY PLAN	
text	
1	Limit (cost=0.42..10.84 rows=20 width=30) (actual time=0.116..0.151 rows=20 loops=1)
2	Buffers: shared hit=20 read=3
3	-> Index Scan using idx_manu_origin_name on manufacturer (cost=0.42..84224.24 rows=161712 width=30) (actual time=0.115..0.147 rows=20 loop..)
4	Index Cond: ((origin)::text = 'USA')::text
5	Buffers: shared hit=20 read=3
6	Planning:
7	Buffers: shared hit=19 read=1
8	Planning Time: 1.908 ms
9	Execution Time: 0.178 ms

Reason: Creating a composite B-tree on (origin, firm\_name) let Postgres use the index as a pointer straight to the origin='USA' range instead of scanning the whole table, and because B-trees support **ordered access**, rows come out already sorted by firm\_name—so the planner switches from a Sequential Scan + sort to an ordered Index Scan that can stop at LIMIT 20.

#### 4. Origin + category filter, sorted by name

```
SELECT firm_name, id
FROM manufacturer
WHERE origin = 'Germany' AND product_category = 'Automotive'
ORDER BY firm_name;
```

Purpose: Germany automotive makers, A→Z.

Index: B-tree covering composite index

```
CREATE INDEX idx_origin_category_name
ON manufacturer (origin, product_category, firm_name);
```

**Before:**

QUERY PLAN
text
Sort (cost=26162.13..26162.14 rows=1 width=23) (actual time=124.653..132.542 rows=0 loops=1)
Sort Key: firm_name
Sort Method: quicksort Memory: 25kB
Buffers: shared hit=15995 read=4056
-> Gather (cost=1000.00..26162.12 rows=1 width=23) (actual time=124.641..132.528 rows=0 loops=1)
Workers Planned: 2
Workers Launched: 2
Buffers: shared hit=15992 read=4056
-> Parallel Seq Scan on manufacturer (cost=0.00..25162.03 rows=1 width=23) (actual time=82.381..82.382 rows=0 loops=1)
Filter: (((origin)::text = 'Germany'::text) AND ((product_category)::text = 'Automotive'::text))
Rows Removed by Filter: 272748
Buffers: shared hit=15992 read=4056
Planning:
Buffers: shared hit=159
Planning Time: 2.794 ms
Execution Time: 132.567 ms

**After:**

QUERY PLAN
text
Index Scan using idx_origin_category_name on manufacturer (cost=0.42..8.45 rows=1 width=23) (actual time=0.071..0.071 rows=0 loops=1)
Index Cond: (((origin)::text = 'Germany'::text) AND ((product_category)::text = 'Automotive'::text))
Buffers: shared read=3
Planning:
Buffers: shared hit=24 read=1
Planning Time: 1.616 ms
Execution Time: 0.086 ms

Reason: Creating a composite B-tree on (origin, product\_category, firm\_name) lines up with both the equality filters and the ORDER BY, letting PostgreSQL use an Index Scan to jump via the index “pointer” straight to the matching, already-sorted key range—so it reads far fewer pages and avoids a table scan and extra sort.

### 5. Top employers in a country (USA)

```
SELECT firm_name, employee_count
FROM manufacturer
WHERE origin = 'USA'
ORDER BY employee_count DESC
LIMIT 20;
```

Purpose: Biggest US employers, top 20 by headcount.  
Index: B-tree composite index

```
CREATE INDEX idx_origin_empdesc_cover
ON manufacturer (origin, employee_count DESC)
INCLUDE (firm_name);
```

*Before:*

QUERY PLAN
text
Limit (cost=27102.67..27105.00 rows=20 width=23) (actual time=124.570..134.119 rows=20 loops=1)
Buffers: shared hit=16084 read=4036
-> Gather Merge (cost=27102.67..42825.77 rows=134760 width=23) (actual time=124.568..134.114 rows=20 loops=1)
Workers Planned: 2
Workers Launched: 2
Buffers: shared hit=16084 read=4036
-> Sort (cost=26102.65..26271.10 rows=67380 width=23) (actual time=81.300..81.301 rows=20 loops=3)
Sort Key: employee_count DESC
Sort Method: top-N heapsort Memory: 27kB
Buffers: shared hit=16084 read=4036
Worker 0: Sort Method: top-N heapsort Memory: 26kB
Worker 1: Sort Method: top-N heapsort Memory: 27kB
-> Parallel Seq Scan on manufacturer (cost=0.00..24309.69 rows=67380 width=23) (actual time=0.030..73.500 rows=5438...
Filter: ((origin)::text = 'USA')::text
Rows Removed by Filter: 218365
Buffers: shared hit=16012 read=4036
Planning:
Buffers: shared hit=129 read=10
Planning Time: 11.373 ms
Execution Time: 134.168 ms

*After:*

QUERY PLAN
text
Limit (cost=0.42..1.31 rows=20 width=23) (actual time=0.473..0.513 rows=20 loops=1)
Buffers: shared hit=1 read=4
-> Index Only Scan using idx_origin_empdesc_cover on manufacturer (cost=0.42..7156.89 rows=161712 width=23) (actual time=0.471..0.509 rows=20 loop...
Index Cond: (origin = 'USA')::text
Heap Fetches: 0
Buffers: shared hit=1 read=4
Planning:
Buffers: shared hit=24 read=1
Planning Time: 1.647 ms
Execution Time: 0.532 ms

Reason: Creating a covering composite B-tree on (origin, employee\_count DESC) INCLUDE (firm\_name) aligns the filter and DESC order, so the optimizer can walk the index in order and stop after the first 20 entries via an Index Only Scan, avoiding a full-table read and separate sort.

		<p><b>6. Operational electronics firms, sorted by name</b></p> <pre> SELECT firm_name, id FROM manufacturer WHERE product_category = 'Electronics' AND is_operational = true ORDER BY firm_name LIMIT 50; </pre> <p>Purpose: Active electronics companies, A→Z, first 50. Index: B-tree Partial, covering, composite</p> <pre> CREATE INDEX idx_cat_oper_name_cover ON manufacturer (product_category, is_operational, firm_name) INCLUDE (id) WHERE is_operational = true; </pre> <p style="text-align: center;"><b>Before:</b></p> <div style="background-color: #f0f0f0; padding: 5px;"> <p><b>QUERY PLAN</b></p> <p>Limit (cost=25870.98..25876.82 rows=50 width=23) (actual time=140.260..152.622 rows=50 loops=1)</p> <p>Buffers: shared hit=16206 read=3916</p> <p>-&gt; Gather Merge (cost=25870.98..29813.66 rows=33792 width=23) (actual time=140.258..152.612 rows=50 loops=1)</p> <p>Workers Planned: 2</p> <p>Workers Launched: 2</p> <p>Buffers: shared hit=16206 read=3916</p> <p>-&gt; Sort (cost=24870.96..24913.20 rows=16896 width=23) (actual time=93.870..93.876 rows=44 loops=3)</p> <p>Sort Key: firm_name</p> <p>Sort Method: top-N heapsort Memory: 30kB</p> <p>Buffers: shared hit=16206 read=3916</p> <p>Worker 0: Sort Method: top-N heapsort Memory: 30kB</p> <p>Worker 1: Sort Method: top-N heapsort Memory: 30kB</p> <p>-&gt; Parallel Seq Scan on manufacturer (cost=0.00..24309.69 rows=16896 width=23) (actual time=0.315..85.843 rows=1356..)</p> <p>Filter: (is_operational AND ((product_category)::text = 'Electronics'::text))</p> <p>Rows Removed by Filter: 259187</p> <p>Buffers: shared hit=16132 read=3916</p> <p>Planning Time: 0.141 ms</p> <p>Execution Time: 152.659 ms</p> </div> <p style="text-align: center;"><b>After:</b></p> <div style="background-color: #f0f0f0; padding: 5px;"> <p><b>QUERY PLAN</b></p> <p>Limit (cost=0.42..3.00 rows=50 width=23) (actual time=0.337..0.366 rows=50 loops=1)</p> <p>Buffers: shared hit=1 read=4</p> <p>-&gt; Index Only Scan using idx_cat_oper_name_cover on manufacturer (cost=0.42..2093.48 rows=40550 width=23) (actual time=0.336..0.362 rows=50 loop...)</p> <p>Index Cond: (product_category = 'Electronics'::text)</p> <p>Heap Fetches: 0</p> <p>Buffers: shared hit=1 read=4</p> <p>Planning:</p> <p>Buffers: shared hit=26 read=1</p> <p>Planning Time: 1.544 ms</p> <p>Execution Time: 0.380 ms</p> </div> <p>Reason: Creating a partial, covering composite B-tree on (product_category, is_operational, firm_name) INCLUDE (id) WHERE is_operational = true aligned the filter and A→Z order, so PostgreSQL used an Index Only Scan over a small, already-sorted subset and stopped at LIMIT 50, avoiding a table scan and extra sort.</p>	
9	Report Quality	<i>Based on quality of documentation for above criteria. No inputs required here.</i>	3/5
10	Demonstration & Interview	<i>Based on assessment during demonstration &amp; interview. No inputs required here.</i>	-